

Topics:

- Neural Networks
- Backpropagation

**CS 4644-DL / 7643-A**  
**ZSOLT KIRA**

- **Assignment 1 out!**

- **Due Feb 5<sup>th</sup>**
- Start now, start now, start now!
- Start now, start now, start now!
- Start now, start now, start now!


- **Piazza**

- Be active!!!
- Extra credit!

- **Office hours**

- [Assignment](#) (@41) and [matrix calculus](#) (@46)


Example with an image with **4 pixels**, and **3 classes** (**cat**/**dog**/**ship**)



Adapted from slides by Fei-Fei Li, Justin Johnson, Serena Yeung, from CS 231n

Example

- We can find the steepest descent direction by computing the **derivative (gradient)**:
$$f'(a) = \lim_{h \rightarrow 0} \frac{f(a + h) - f(a)}{h}$$
- Steepest descent direction is the **negative gradient**
- **Intuitively:** Measures how the function changes as the argument  $a$  changes by a small step size
  - As step size goes to zero
- **In Machine Learning:** Want to know how the **loss function** changes **as weights** are varied
  - Can consider each parameter separately by taking **partial derivative** of loss function with respect to that parameter



This idea can be turned into an **algorithm (gradient descent)**

- Choose a model:  $f(x, W) = Wx$
- Choose loss function:  $L_i = |y - Wx_i|^2$
- Calculate partial derivative for each parameter:  $\frac{\partial L}{\partial w_i}$
- Update the parameters:  $w_i = w_i - \frac{\partial L}{\partial w_i}$
- Add learning rate to prevent too big of a step:  $w_i = w_i - \alpha \frac{\partial L}{\partial w_i}$
- Repeat (from Step 3)

Often, we only compute the gradients across a small subset of data

● Full Batch Gradient Descent 
$$L = \frac{1}{N} \sum L(f(x_i, W), y_i)$$

● Mini-Batch Gradient Descent 
$$L = \frac{1}{M} \sum L(f(x_i, W), y_i)$$


● Where M is a *subset* of data

● We iterate over mini-batches:

● Get mini-batch, compute loss, compute derivatives, and take a set

## Mini-Batch Gradient Descent

<http://demonstrations.wolfram.com/VisualizingTheGradientVector/>



**Gradient Descent**

For some functions, we can analytically derive the partial derivative

## Example:

### Function

$$f(\mathbf{w}, \mathbf{x}_i) = \mathbf{w}^T \mathbf{x}_i$$

(Assume  $\mathbf{w}$  and  $\mathbf{x}_i$  are column vectors, so same as  $\mathbf{w} \cdot \mathbf{x}_i$ )

### Loss

$$(\mathbf{y}_i - \mathbf{w}^T \mathbf{x}_i)^2$$

**Dataset:**  $N$  examples (indexed by  $k$ )

### Update Rule

$$\mathbf{w}_j \leftarrow \mathbf{w}_j + 2\eta \sum_{k=1}^N \delta_k \mathbf{x}_{kj}$$

## Derivation of Update Rule

$$L = \sum_{k=1}^N (\mathbf{y}_k - \mathbf{w}^T \mathbf{x}_k)^2$$

Gradient descent tells us we should update  $\mathbf{w}$  as follows to minimize  $L$ :

$$\mathbf{w}_j \leftarrow \mathbf{w}_j - \eta \frac{\partial L}{\partial \mathbf{w}_j}$$

So what's  $\frac{\partial L}{\partial \mathbf{w}_j}$ ?

$$\begin{aligned} \frac{\partial L}{\partial \mathbf{w}_j} &= \sum_{k=1}^N \frac{\partial}{\partial \mathbf{w}_j} (\mathbf{y}_k - \mathbf{w}^T \mathbf{x}_k)^2 \\ &= \sum_{k=1}^N 2(\mathbf{y}_k - \mathbf{w}^T \mathbf{x}_k) \frac{\partial}{\partial \mathbf{w}_j} (\mathbf{y}_k - \mathbf{w}^T \mathbf{x}_k) \\ &= -2 \sum_{k=1}^N \delta_k \frac{\partial}{\partial \mathbf{w}_j} \mathbf{w}^T \mathbf{x}_k \quad \boxed{\dots \text{where } \dots} \\ &\quad \delta_k = \mathbf{y}_k - \mathbf{w}^T \mathbf{x}_k \\ &= -2 \sum_{k=1}^N \delta_k \frac{\partial}{\partial \mathbf{w}_j} \sum_{i=1}^m \mathbf{w}_i \mathbf{x}_{ki} \\ &= -2 \sum_{k=1}^N \delta_k \mathbf{x}_{kj} \end{aligned}$$



If we add a **non-linearity (sigmoid)**, derivation is more complex

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$


First, one can derive that:  $\sigma'(x) = \sigma(x)(1 - \sigma(x))$

$$\mathbf{f}(\mathbf{x}) = \sigma\left(\sum_k w_k x_k\right)$$

$$L = \sum_i \left( y_i - \sigma\left(\sum_k w_k x_{ik}\right) \right)^2$$

$$\begin{aligned} \frac{\partial L}{\partial w_j} &= \sum_i 2 \left( y_i - \sigma\left(\sum_k w_k x_{ik}\right) \right) \left( -\frac{\partial}{\partial w_j} \sigma\left(\sum_k w_k x_{ik}\right) \right) \\ &= \sum_i -2 \left( y_i - \sigma\left(\sum_k w_k x_{ik}\right) \right) \sigma'\left(\sum_k w_k x_{ik}\right) \frac{\partial}{\partial w_j} \sum_k w_k x_{ik} \\ &= \sum_i -2 \delta_i \sigma(d_i) (1 - \sigma(d_i)) x_{ij} \end{aligned}$$

where  $\delta_i = y_i - \mathbf{f}(x_i)$        $d_i = \sum_k w_k x_{ik}$



The sigmoid perception update rule:

$$w_j \leftarrow w_j + 2\eta \sum_{k=1}^N \delta_i \sigma_i (1 - \sigma_i) x_{ij}$$

where  $\sigma_i = \sigma\left(\sum_{j=1}^m w_j x_{ij}\right)$


$$\delta_i = y_i - \sigma_i$$

## Adding a Non-Linear Function

A **linear classifier** can be broken down into:

- ◆ Input
- ◆ A function of the input
- ◆ A loss function

It's all just one function that can be **decomposed** into building blocks



What Does a Linear Classifier Consist of?

The same two-layered neural network corresponds to adding another weight matrix

- We will prefer the linear algebra view, but use some terminology from neural networks (& biology)




Figure adapted from slides by Fei-Fei Li, Justin Johnson, Serena Yeung, CS 231n



## The Linear Algebra View

**Large (deep) networks** can be built by adding more and more layers

Three-layered neural networks can represent **any function**

- The number of nodes could grow unreasonably (exponential or worse) with respect to the complexity of the function

We will show them **without edges**:



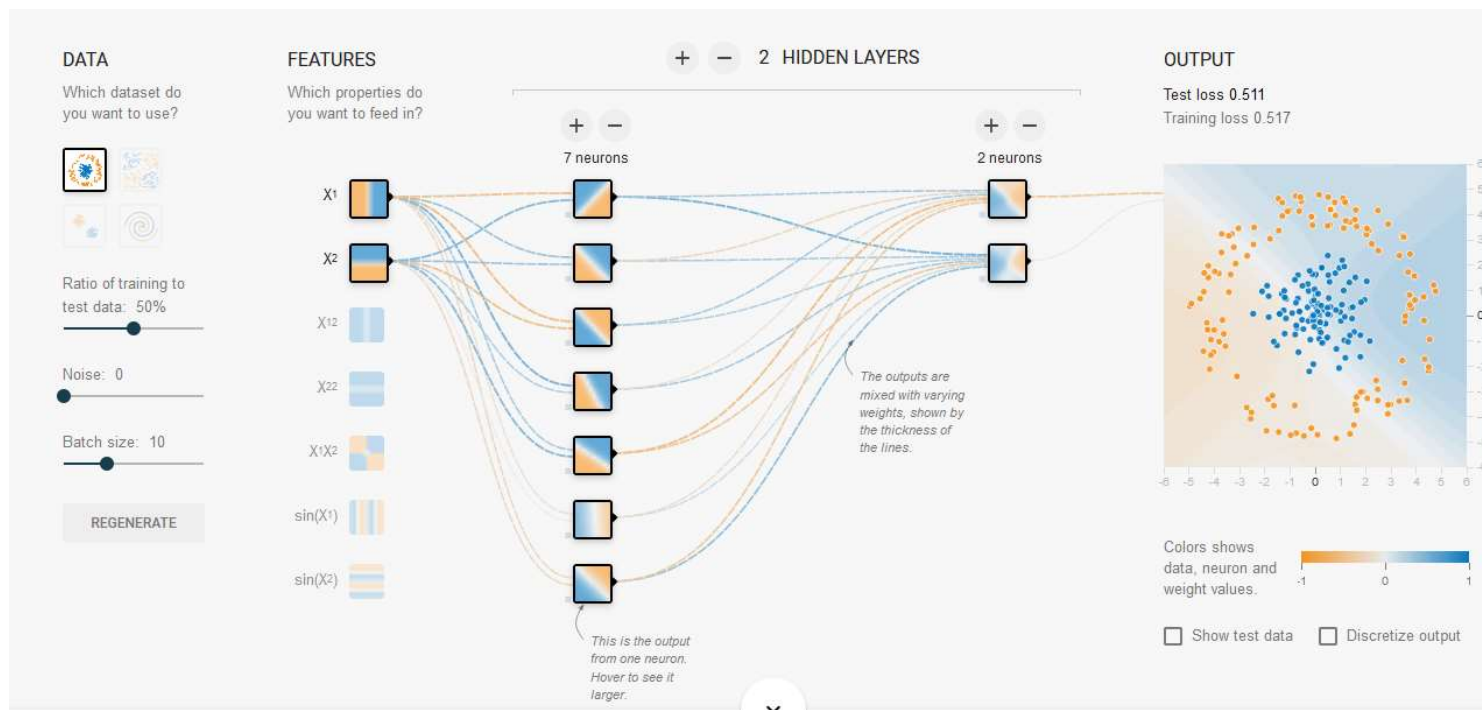
$$f(x, W_1, W_2, W_3) = \sigma(W_2 \sigma(W_1 x))$$

Figure adapted from slides by Fei-Fei Li, Justin Johnson, Serena Yeung, CS 231n

**Adding More Layers!**

# Demo

- <http://playground.tensorflow.org>



# Computation Graphs


Functions can be made **arbitrarily complex** (subject to memory and computational limits), e.g.:

$$f(x, W) = \sigma(W_5 \sigma(W_4 \sigma(W_3 \sigma(W_2 \sigma(W_1 x))))$$

We can use **any type of differentiable function (layer)** we want!

◆ At the end, **add the loss function**

Composition can have **some structure**




**Adding Even More Layers**

The world is **compositional**!

We want our **model** to reflect this

Empirical and theoretical evidence that it makes **learning complex functions easier**

Note that **prior state of art engineered features** often had this compositionality as well




*Adapted from figure by Marc'Aurelio Ranzato, Yann LeCun*

◆ **Pixels -> edges -> object parts -> objects**

**Compositionality**




- We are learning **complex models** with significant amount of parameters (millions or billions)
- How do we compute the gradients of the **loss** (at the end) with respect to **internal** parameters?
- Intuitively, want to understand how **small changes** in weight deep inside **are propagated** to affect the **loss function** at the end



## Computing Gradients in Complex Function

Given a library of simple functions



Compose into a  
→  
complicate function

$$-\log\left(\frac{1}{1 + e^{-w \cdot x}}\right)$$



*Adapted from slides by: Marc'Aurelio Ranzato, Yann LeCun*


## Decomposing a Function

To develop a general algorithm for this, we will view the function as a **computation graph**

Graph can be any **directed acyclic graph (DAG)**

- Modules must be differentiable to support gradient computations for gradient descent



A **training algorithm** will then process this graph, **one module at a time**



*Adapted from figure by Marc'Aurelio Ranzato, Yann LeCun*



# Directed Acyclic Graphs (DAGs)

- Exactly what the name suggests
  - Directed edges
  - No (directed) cycles
  - Underlying undirected cycles okay




# Directed Acyclic Graphs (DAGs)


- Concept
  - Topological Ordering



# Directed Acyclic Graphs (DAGs)




$$f(x_1, x_2) = \ln(x_1) + x_1x_2 - \sin(x_2)$$



Example

$$-\log\left(\frac{1}{1 + e^{-w \cdot x}}\right)$$



*Adapted from figure by Marc'Aurelio Ranzato, Yann LeCun*



# Backpropagation

## Step 1: Compute Loss on Mini-Batch: Forward Pass




*Adapted from figure by Marc'Aurelio Ranzato, Yann LeCun*

## Step 1: Compute Loss on Mini-Batch: Forward Pass



*Adapted from figure by Marc'Aurelio Ranzato, Yann LeCun*

## Step 1: Compute Loss on Mini-Batch: Forward Pass




Note that we must store the **intermediate outputs of all layers!**

- ◆ This is because we will need them to **compute the gradients** (the gradient equations will have terms with the output values in them)

*Adapted from figure by Marc'Aurelio Ranzato, Yann LeCun*

**Step 1:** Compute Loss on Mini-Batch: **Forward Pass**


**Step 2:** Compute Gradients wrt parameters: **Backward Pass**



*Adapted from figure by Marc'Aurelio Ranzato, Yann LeCun*

**Step 1:** Compute Loss on Mini-Batch: **Forward Pass**

**Step 2:** Compute Gradients wrt parameters: **Backward Pass**



*Adapted from figure by Marc'Aurelio Ranzato, Yann LeCun*


**Step 1:** Compute Loss on Mini-Batch: **Forward Pass**

**Step 2:** Compute Gradients wrt parameters: **Backward Pass**



*Adapted from figure by Marc'Aurelio Ranzato, Yann LeCun*

- We want to compute:  $\left\{ \frac{\partial L}{\partial h^{\ell-1}}, \frac{\partial L}{\partial W} \right\}$



- We will use the *chain rule* to do this:

Chain Rule:  $\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \cdot \frac{\partial y}{\partial x}$


## Computing the Gradients of Loss



**Step 1:** Compute Loss on Mini-Batch: **Forward Pass**

**Step 2:** Compute Gradients wrt parameters: **Backward Pass**

**Step 3:** Use **gradient** to update **all parameters** at the end



$$w_i = w_i - \alpha \frac{\partial L}{\partial w_i}$$

**Backpropagation is the application of gradient descent to a computation graph via the chain rule!**



*Adapted from figure by Marc'Aurelio Ranzato, Yann LeCun*


Given this computation graph, the training algorithm will:

- Calculate the current model's outputs (called the **forward pass**)
- Calculate the gradients for each module (called the **backward pass**)

Backward pass is a recursive algorithm that:

- Starts at **loss function** where we know how to calculate the gradients
- Progresses back through the modules
- Ends in the **input layer** where we do not need gradients (no parameters)


This algorithm is called **backpropagation**



*Adapted from figure by Marc'Aurelio Ranzato, Yann LeCun*

In the **backward pass**, we seek to calculate the gradients of the loss with respect to the module's parameters

- Assume that we have the gradient of the loss with respect to the **module's outputs** (given to us by upstream module)
- We will also pass the gradient of the loss with respect to the **module's inputs**
  - This is not required for update the module's weights, but passes the gradients back to the previous module



### Problem:

- We can compute local gradients:  $\left\{ \frac{\partial h^{\ell}}{\partial h^{\ell-1}}, \frac{\partial h^{\ell}}{\partial W} \right\}$
- We are given:  $\frac{\partial L}{\partial h^{\ell}}$
- Compute:  $\left\{ \frac{\partial L}{\partial h^{\ell-1}}, \frac{\partial L}{\partial W} \right\}$

*Adapted from figure by Marc'Aurelio Ranzato, Yann LeCun*

## Backward Pass Computations

- We can compute **local gradients**:  $\{\frac{\partial h^\ell}{\partial h^{\ell-1}}, \frac{\partial h^\ell}{\partial W}\}$
- This is just the **derivative of our function** with respect to its parameters and inputs!


**Example:** If  $h^\ell = Wh^{\ell-1}$

$$\text{then } \frac{\partial h^\ell}{\partial h^{\ell-1}} = W$$

$$\text{and } \frac{\partial h^\ell}{\partial w_i} = h^{\ell-1,T} \quad \begin{array}{l} \text{(a sparse matrix with} \\ \text{in the } i\text{-th row} \end{array}$$

## Computing the Local Gradients: Example

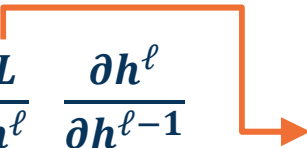
- We want to compute:  $\left\{ \frac{\partial L}{\partial h^{\ell-1}}, \frac{\partial L}{\partial W} \right\}$




- We will use the *chain rule* to do this:

Chain Rule:  $\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \cdot \frac{\partial y}{\partial x}$

- We will use the **chain rule** to compute:  $\{\frac{\partial L}{\partial h^{\ell-1}}, \frac{\partial L}{\partial W}\}$

- **Gradient of loss w.r.t. inputs:**  $\frac{\partial L}{\partial h^{\ell-1}} = \frac{\partial L}{\partial h^{\ell}} \frac{\partial h^{\ell}}{\partial h^{\ell-1}}$   Given by upstream module (**upstream gradient**)

- **Gradient of loss w.r.t. weights:**  $\frac{\partial L}{\partial W} = \frac{\partial L}{\partial h^{\ell}} \frac{\partial h^{\ell}}{\partial W}$




*Adapted from figure by Marc'Aurelio Ranzato, Yann LeCun*

# Backpropagation: a simple example

$$f(x, y, z) = (x + y)z$$

# Backpropagation: a simple example

$$f(x, y, z) = (x + y)z$$






# Backpropagation: a simple example

$$f(x, y, z) = (x + y)z$$


e.g.  $x = -2$ ,  $y = 5$ ,  $z = -4$



# Backpropagation: a simple example

$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$




Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$

# Backpropagation: a simple example

$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$



Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$

# Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



# Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial f}$$

# Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial f}$$

# Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial z}$$

A pink arrow points from this box to the  $z$  input of the multiplication node in the computational graph above.

# Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial z}$$



# Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial q}$$

# Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial q}$$

# Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Chain rule:

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial y}$$

Upstream  
gradient

Local  
gradient

$$\frac{\partial f}{\partial y}$$

# Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Chain rule:

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial y}$$

Upstream  
gradient

Local  
gradient

$$\frac{\partial f}{\partial y}$$

# Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Chain rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$

Upstream  
gradient

Local  
gradient

$$\frac{\partial f}{\partial x}$$

# Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want:  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Chain rule:


$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$

Upstream  
gradient


Local  
gradient

$$\frac{\partial f}{\partial x}$$

# Backpropagation: a simple example




# Backpropagation: a simple example






# Patterns in backward flow




# Patterns in backward flow

Q: What is an **add** gate?



# Patterns in backward flow


**add** gate: gradient distributor



# Patterns in backward flow

**add** gate: gradient distributor


Q: What is a **max** gate?



# Patterns in backward flow

**add** gate: gradient distributor

**max** gate: gradient router




# Patterns in backward flow

**add** gate: gradient distributor

**max** gate: gradient router

Q: What is a **mul** gate?




# Patterns in backward flow


**add** gate: gradient distributor

**max** gate: gradient router

**mul** gate: gradient switcher




# Gradients add at branches






# Duality in Fprop and Bprop



# Deep Learning = Differentiable Programming

- Computation = Graph
  - Input = Data + Parameters
  - Output = Loss
  - Scheduling = Topological ordering
- What do we need to do?
  - Generic code for representing the graph of modules
  - Specify modules (both forward and backward function)


# Modularized implementation: forward / backward API



Graph (or Net) object (*rough psuedo code*)

```
class ComputationalGraph(object):  
    #...  
    def forward(inputs):  
        # 1. [pass inputs to input gates...]  
        # 2. forward the computational graph:  
        for gate in self.graph.nodes_topologically_sorted():  
            gate.forward()  
        return loss # the final gate in the graph outputs the loss  
    def backward():  
        for gate in reversed(self.graph.nodes_topologically_sorted()):  
            gate.backward() # little piece of backprop (chain rule applied)  
        return inputs_gradients
```

## Modularized implementation: forward / backward API



(x,y,z are scalars)

```
class MultiplyGate(object):  
    def forward(x,y):  
        z = x*y  
        return z  
    def backward(dz):  
        # dx = ... #todo  
        # dy = ... #todo  
        return [dx, dy]
```


$$\frac{\partial L}{\partial z}$$

Arrow pointing to the `backward(dz)` parameter in the code block.

$$\frac{\partial L}{\partial x}$$

Arrow pointing to the `dx` element in the `return [dx, dy]` statement in the code block.

## Modularized implementation: forward / backward API



(x,y,z are scalars)

```
class MultiplyGate(object):  
    def forward(x,y):  
        z = x*y  
        self.x = x # must keep these around!  
        self.y = y  
        return z  
    def backward(dz):  
        dx = self.y * dz # [dz/dx * dL/dz]  
        dy = self.x * dz # [dz/dy * dL/dz]  
        return [dx, dy]
```

# Example: Caffe layers

Branch: master - caffe / src / caffe / layers /			Create new file	Upload files	Find file	History
sheelamer committed on GitHub Merge pull request #4630 from BGGens/load_hdfs_fix [ci]			Latest commit e687a71 21 days ago			
absval_layer.cpp	dismantle layer headers	a year ago				
absval_layer.cu	dismantle layer headers	a year ago				
accuracy_layer.cpp	dismantle layer headers	a year ago				
argmax_layer.cpp	dismantle layer headers	a year ago				
base_conv_layer.cpp	enable dilated deconvolution	a year ago				
base_data_layer.cpp	Using default from proto for prefetch	3 months ago				
base_data_layer.cu	Switched multi-GPU to NCCL	3 months ago				
batch_norm_layer.cpp	Add missing spaces besides equal signs in batch_norm_layer.cpp	4 months ago				
batch_norm_layer.cu	dismantle layer headers	a year ago				
batch_reindex_layer.cpp	dismantle layer headers	a year ago				
batch_reindex_layer.cu	dismantle layer headers	a year ago				
bias_layer.cpp	Remove incorrect cast of gemm int arg to Dtype in BiasLayer	a year ago				
bias_layer.cu	Separation and generalization of ChannelwiseAffineLayer into BiasLayer	a year ago				
bnll_layer.cpp	dismantle layer headers	a year ago				
bnll_layer.cu	dismantle layer headers	a year ago				
concat_layer.cpp	dismantle layer headers	a year ago				
concat_layer.cu	dismantle layer headers	a year ago				
contrastive_loss_layer.cpp	dismantle layer headers	a year ago				
contrastive_loss_layer.cu	dismantle layer headers	a year ago				
conv_layer.cpp	add support for 2D dilated convolution	a year ago				
conv_layer.cu	dismantle layer headers	a year ago				
crop_layer.cpp	remove redundant operations in Crop layer (#5138)	2 months ago				
crop_layer.cu	remove redundant operations in Crop layer (#5138)	2 months ago				
cudnn_conv_layer.cpp	dismantle layer headers	a year ago				
cudnn_conv_layer.cu	Add cuDNN v5 support, drop cuDNN v3 support	11 months ago				

Caffe is licensed under [BSD 2-Clause](#)

cudnn_lcn_layer.cpp	dismantle layer headers	a year ago
cudnn_lcn_layer.cu	dismantle layer headers	a year ago
cudnn_lrn_layer.cpp	dismantle layer headers	a year ago
cudnn_lrn_layer.cu	dismantle layer headers	a year ago
cudnn_pooling_layer.cpp	dismantle layer headers	a year ago
cudnn_pooling_layer.cu	dismantle layer headers	a year ago
cudnn_relu_layer.cpp	Add cuDNN v5 support, drop cuDNN v3 support	11 months ago
cudnn_relu_layer.cu	Add cuDNN v5 support, drop cuDNN v3 support	11 months ago
cudnn_sigmoid_layer.cpp	Add cuDNN v5 support, drop cuDNN v3 support	11 months ago
cudnn_sigmoid_layer.cu	Add cuDNN v5 support, drop cuDNN v3 support	11 months ago
cudnn_softmax_layer.cpp	dismantle layer headers	a year ago
cudnn_softmax_layer.cu	dismantle layer headers	a year ago
cudnn_tanh_layer.cpp	Add cuDNN v5 support, drop cuDNN v3 support	11 months ago
cudnn_tanh_layer.cu	Add cuDNN v5 support, drop cuDNN v3 support	11 months ago
data_layer.cpp	Switched multi-GPU to NCCL	3 months ago
deconv_layer.cpp	enable dilated deconvolution	a year ago
deconv_layer.cu	dismantle layer headers	a year ago
dropout_layer.cpp	supporting N-D Blobs in Dropout layer Reshape	a year ago
dropout_layer.cu	dismantle layer headers	a year ago
dummy_data_layer.cpp	dismantle layer headers	a year ago
eltwise_layer.cpp	dismantle layer headers	a year ago
eltwise_layer.cu	dismantle layer headers	a year ago
elu_layer.cpp	ELU layer with basic tests	a year ago
elu_layer.cu	ELU layer with basic tests	a year ago
embed_layer.cpp	dismantle layer headers	a year ago
embed_layer.cu	dismantle layer headers	a year ago
euclidean_loss_layer.cpp	dismantle layer headers	a year ago
euclidean_loss_layer.cu	dismantle layer headers	a year ago
exp_layer.cpp	Solving issue with exp layer with base e	a year ago
exp_layer.cu	dismantle layer headers	a year ago

# Caffe Sigmoid Layer

```

1 #include <cmath>
2 #include <vector>
3
4 #include "caffe/layers/sigmoid_layer.hpp"
5
6 namespace caffe {
7
8 template <typename Dtype>
9 inline Dtype sigmoid(Dtype x) {
10     return 1. / (1. + exp(-x));
11 }
12
13 template <typename Dtype>
14 void SigmoidLayer<Dtype>::Forward_cpu(const vector<Blob<Dtype>>*& bottom,
15     const vector<Blob<Dtype>>*& top) {
16     const Dtype* bottom_data = bottom[0]->cpu_data();
17     Dtype* top_data = top[0]->mutable_cpu_data();
18     const int count = bottom[0]->count();
19     for (int i = 0; i < count; ++i) {
20         top_data[i] = sigmoid(bottom_data[i]);
21     }
22 }
23
24 template <typename Dtype>
25 void SigmoidLayer<Dtype>::Backward_cpu(const vector<Blob<Dtype>>*& top,
26     const vector<bool>& propagate_down,
27     const vector<Blob<Dtype>>*& bottom) {
28     if (propagate_down[0]) {
29         const Dtype* top_data = top[0]->cpu_data();
30         const Dtype* top_diff = top[0]->cpu_diff();
31         Dtype* bottom_diff = bottom[0]->mutable_cpu_diff();
32         const int count = bottom[0]->count();
33         for (int i = 0; i < count; ++i) {
34             const Dtype sigmoid_x = top_data[i];
35             bottom_diff[i] = top_diff[i] * sigmoid_x * (1. - sigmoid_x);
36         }
37     }
38 }
39
40 #ifndef CPU_ONLY
41 STUB_GPU(SigmoidLayer);
42 #endif
43
44 INSTANTIATE_CLASS(SigmoidLayer);
45
46 } // namespace caffe

```

[Caffe](#) is licensed under [BSD 2-Clause](#)

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

$$(1 - \sigma(x)) \sigma(x) * \text{top\_diff} \text{ (chain rule)}$$

Batches of data are **matrices** or **tensors** (multi-dimensional matrices)

### Examples:

- Each instance is a vector of size  $m$ , our batch is of size  $[B \times m]$
- Each instance is a matrix (e.g. grayscale image) of size  $W \times H$ , our batch is  $[B \times W \times H]$
- Each instance is a multi-channel matrix (e.g. color image with R,B,G channels) of size  $C \times W \times H$ , our batch is  $[B \times C \times W \times H]$

### Jacobians become tensors which is complicated

- Instead, flatten input to a vector and get a vector of derivatives!
- In practice, figure out Jacobians for simpler items (scalars, vectors), figure out pattern, and *slice* or index appropriate elements to create Jacobians


$$\begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1n} \\ x_{21} & x_{22} & \cdots & x_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \cdots & x_{nn} \end{bmatrix}$$

Flatten



$$\begin{bmatrix} x_{11} \\ x_{12} \\ \vdots \\ x_{21} \\ x_{22} \\ \vdots \\ x_{n1} \\ \vdots \\ x_{nn} \end{bmatrix}$$






$$h^{\ell} = W h^{\ell-1}$$

$$\begin{array}{c}
 \left[ \begin{array}{c} \\ \\ \end{array} \right] \quad \left[ \begin{array}{c} \leftarrow w_i^T \rightarrow \\ \\ \end{array} \right] \quad \left[ \begin{array}{c} \\ \\ \end{array} \right] \\
 |h^{\ell}| \times 1 \quad |h^{\ell}| \times |h^{\ell-1}| \quad |h^{\ell-1}| \times 1
 \end{array}$$

## Fully Connected (FC) Layer: Forward Function



Note doing this on full  $W$  matrix would result in Jacobian tensor!

But it is *sparse* – each output only affected by corresponding weight row

$$\frac{\partial \mathbf{h}^{\ell}}{\partial \mathbf{h}^{\ell-1}} = W$$

$$\frac{\partial L}{\partial \mathbf{h}^{\ell-1}} = \frac{\partial L}{\partial \mathbf{h}^{\ell}} \frac{\partial \mathbf{h}^{\ell}}{\partial \mathbf{h}^{\ell-1}}$$

$$\begin{bmatrix} & \end{bmatrix} \begin{bmatrix} & \end{bmatrix} \begin{bmatrix} \\ \\ \\ \end{bmatrix}$$

$$\frac{\partial L}{\partial w_i} = \frac{\partial L}{\partial \mathbf{h}^{\ell}} \frac{\partial \mathbf{h}^{\ell}}{\partial w_i}$$

$$\begin{bmatrix} & \end{bmatrix} \begin{bmatrix} & \end{bmatrix} \begin{bmatrix} \leftarrow 0 \rightarrow \\ \leftarrow \frac{\partial \mathbf{h}^{\ell}}{\partial w_i} \rightarrow \\ \leftarrow 0 \rightarrow \end{bmatrix}$$

$$\frac{\partial \mathbf{h}^{\ell}}{\partial w_i} = \mathbf{h}^{(\ell-1),T}$$

$$\begin{matrix} 1 \times |\mathbf{h}^{\ell-1}| & 1 \times |\mathbf{h}^{\ell}| & |\mathbf{h}^{\ell}| \times |\mathbf{h}^{\ell-1}| & 1 \times |\mathbf{h}^{\ell-1}| & 1 \times |\mathbf{h}^{\ell}| & |\mathbf{h}^{\ell}| \times |\mathbf{h}^{\ell-1}| \end{matrix}$$

## Fully Connected (FC) Layer